

Final, Spring 2015

This test has 14 questions worth a total of 60 points. The exam is closed book, except that you are allowed to use three (front-and-back) handwritten pages of notes. No calculators or other electronic devices are permitted. Give your answers and show your work in the space provided. **Write the statement out below in the blank provided and sign. You may do this before the exam begins.**

"I have neither given nor received any assistance in the taking of this exam."

	Points		Points
0	1	8	4
1	6	9	4.5
2	3	10	7
3	4.5	11	0
4	7.5	12	8
5	2	13	5
6	2.5		
7	5		

Signature: _____

Your Name: _____

Your Student ID: _____

Three-letter Login ID: _____

Login of Person to Left: _____

Login of Person to Right: _____

Exam Room: _____

Tips:

- There may be partial credit for incomplete answers. Write as much of the solution as you can, but bear in mind that we may deduct points if your answers are much more complicated than necessary.
- There are a lot of problems on this exam. Work through the ones with which you are comfortable first. Do not get overly captivated by interesting design issues or complex corner cases you're not sure about.
- Not all information provided in a problem may be useful.
- All code that we've provided on this exam should compile. All code has been compiled and executed before printing, but in the unlikely event that we do happen to catch any bugs in the exam, we'll announce a fix. The correct answer is not 'does not compile.'
- Don't panic! Recall that we shoot for around a 60% median. You should not expect to be able to do all of the problems on this exam.
- If you're feeling anxious and need to take a break, go for it. Just let a TA know, and leave your test and electronic devices behind.


Optional. Mark along the line to show your feelings
on the spectrum between ☹ and ☺.

Before exam: [☹_____☺].
After exam: [☹_____☺].

0. Another point. (1 points). Write your name, login, and ID on the front page. Write your exam room. Write the IDs of your neighbors. Write the given statement and sign. Write your login in the corner of every page. ☺

1. Basic Operations (6 Points).

- a. For the graph below, use Kruskal's algorithm to find the MST. The number on each edge is the weight, and the letter is a unique label you should use in your answer to specify that edge. Provide the edges **in the order they'd be found** by Kruskal's algorithm. Break any ties using the alphabetical label. Use the blanks below. You may not need all blanks. **Give your answer in terms of the alphabetical labels**, e.g. if Kruskal's starts with the edge between vertices 3 and 4, you would write H in the first blank.



- b. Repeat part a, but using Prim's algorithm, starting from vertex #3. **As before, give your answer in the order added and in terms of the alphabetical labels.** You may not need all of the blanks.

- c. Suppose we have the array [4, 1, 6, 2, 3, 7, 3, 4]. Give a valid partitioning of this array, using the leftmost item as a pivot. Any partitioning scheme is fine.

- d. Give the array that results if we use transform the array [1, 2, 4, 9, 3, 7, 5] into a max heap. Use the sinking based heapify process that we described in class (possibly helpful reminder: this process also happens to be the first step of heap sort).

Login: _____

2. Basic Operations Turbo Edition (3 Points).

- a. Draw a binary tree that has a pre-order and in-order traversal of A, B, C, D, E.
- b. Draw a **directed** graph whose DFS pre-order traversal is A, B, C, D, E, and whose DFS post-order traversal is C, B, D, E, A. Assume that ties are broken alphabetically.

3. Best and Worst (4.5 Points). Give bests and worsts as requested below.

- a. Give an input that results in worst case runtime for insertion sort. Give your answer as an array of 5 integers, written in the blanks below.

- b. Given an initially empty BST, give an insertion order for A, B, C, D, E, F, G that minimizes height. Assume we are not performing any balancing operations.

- c. Suppose we want to build an array based stack that supports push and pop. In class, we learned how appropriate resizing rules can empower data structures with excellent amortized runtime. Suppose our stack has the following two rules:

- Before each push, if the array **is full, double** the size before pushing.
- After each pop, if the array is less than **half full, halve** the size of the array.

What is the worst-case amortized runtime for a sequence of pushes and pops? Give your answer in $\Theta(\cdot)$ notation in terms of N , the maximum size of the array that is reached during the stack's lifetime.

4. Comparative Data Structures and Algorithms (7.5 points).

- a. What is the size of the largest binary min heap that is also a valid BST? Draw an example assuming all keys are letters from the English alphabet. **Assume the heap has no duplicate elements.**

- b. What is the size of the largest binary max heap that is also a valid BST? Draw an example assuming all keys are letters from the English alphabet. **Assume the heap has no duplicate elements.**

- c. What is the size of the largest binary min heap in which the fifth largest element is a child of the root? You do not need to draw an example. **Assume the heap has no duplicate elements.**

- d. Give an example of when you'd use Quicksort over Mergesort.

- e. Give an example of when you'd use a Trie-based map instead of a HashMap.

- f. Why did we bother introducing heaps for implementing priority queues instead of just using left-leaning red-black binary search trees?

Hier ist eine keine Tausendfüßler Zone.

Login: _____

5. Syntax Mastery (2 points). Give the output of `main`. You may not need all lines provided.

```
public class Sklarp implements Iterable<Character>, Iterator<Character> {
    public char[] contents;
    public char magicCharacter;
    public int k;

    public Sklarp(char[] s, Character c) {
        contents = s;
        magicCharacter = c;
        k = 0;
    }

    public Iterator<Character> iterator() {
        return this;
    }

    public boolean hasNext() {
        return k < contents.length;
    }

    public Character next() {
        if (k % 3 == 0) {
            contents[k] = magicCharacter;
        }
        char returnChar = contents[k];
        k += 1;
        return returnChar;
    }

    public void remove() { throw new UnsupportedOperationException(); }

    public static void main(String[] args) {
        Sklarp g = new Sklarp("Zeoi dei".toCharArray(), 'r');
        for (Character c : g) {
            System.out.print(c);
        }
        System.out.println();
        for (Character c : g) {
            System.out.print(c);
        }
        System.out.println();    }    }
```

6. Runtime Analysis (2.5 points). For each of the pieces of code below, give the **worst case** runtime in $\Theta(\cdot)$ notation as a function of N . Your answer should be as simple as possible (i.e. avoid unnecessary constants, lower order terms, etc.). If the worst case is an infinite loop, write an infinity symbol in the blank. Assume there is no limit to the size of an int (otherwise technically they're all constant).

_____ `public static void f1(int N) {
 int sum = 0;
 for (int i = N; i > 0; i -= 1) {
 for (int j = 0; j < i; j += 1) {
 sum += 1;
 }
 }
 }`

_____ `public static void f2(int N) {
 int sum = 0;
 for (int i = 1; i <= N; i = i*2) {
 for (int j = 0; j < i; j += 1) {
 sum += 1;
 }
 }
 }`

_____ `public static void f3(int[] a) {
 if (a.length == 0) { return; }
 int N = a.length;
 int[] newA = new int[N-1];
 for (int i = 0; i < newA.length; i += 1) {
 newA[i] = a[i];
 }
 f3(newA);
 }`

_____ `public static void f4(int N) {
 int x = N;
 while (x > 0) {
 x = x >>> 1;
 }
 }`

_____ `public static void f5(int N) {
 f1(N); // If you left one or more of the answers
 f2(N); // above blank, give your answer to f5 in
 f3(new int[N]); // terms of your non-blank answers.
 f4(N); }`

Login: _____

7. Sorting Mechanics (5 points).

- a. Below, the leftmost column is an array of strings to be sorted. The column to the far right gives the strings in sorted order. Each of the remaining columns gives the contents of the array during some intermediate step of one of the algorithms listed below. Match each column with its corresponding algorithm. Use every answer exactly once. Write your answers in the blanks provided.

4873	1876	1874	1626	9573	2212	1626
1874	1874	1626	1874	7121	8917	1874
8917	2212	1876	1876	9132	7121	1876
1626	1626	1897	4873	6973	1626	1897
4982	3492	2212	4982	4982	9132	2212
9132	1897	3492	8917	8917	6152	3492
9573	4873	4873	9132	6152	4873	4873
1876	9573	4982	9573	1876	9573	4982
6973	6973	6973	1897	1626	6973	6152
1897	9132	6152	3492	1897	1874	6973
9587	9587	7121	6973	1874	1876	7121
3492	4982	8917	9587	3492	9877	8917
9877	9877	9132	2212	4873	4982	9132
2212	8917	9573	6152	2212	9587	9573
6152	6152	9587	7121	9587	3492	9587
7121	7121	9877	9877	9877	1897	9877
----	----	----	----	----	----	----
__1__	_____	_____	_____	_____	_____	__7__

1: Unsorted 2: Merge 3: Quick 4: Heap 5: LSD 6: MSD 7: Sorted

Notes:

- Quicksort: Non-random and uses topmost item as pivot. We have deliberately omitted information about the partitioning strategy.
- Mergesort: Recursive implementation (a.k.a. top-down for textbook readers).
- Heapsort, LSD Radix Sort, and MSD Radix Sort: As described in class.

8. Choosing a Sort (4 Points).

You're an imperial engineer doing some QA on a recent batch of droids that were produced. You have a supposedly sorted array of n Droid objects that implement Comparable. However, when looking through your array, you realize these aren't the droids you're looking for! Your supervisor tells you that the machine malfunctioned, and it's made at most k mistakes: There are no more than k inversions, where **we define an inversion as a pair of droids that is not in the right order.**¹

- a. State the most efficient sorting algorithm and the big O runtime (giving the tightest bound with no unnecessary constants or lower order terms) to sort the droids if:

$$k = O(n)$$

Algorithm: _____

Runtime: _____

$$k = O(n^2)$$

Algorithm: _____

Runtime: _____

$$k = \log(n)$$

Algorithm: _____

Runtime: _____

- b. Two weeks later, you're given another batch of Droids that are supposed to be sorted on a 32-bit int ID, an instance variable of Droid. The sorting machine hasn't been fixed yet and again has made at most k mistakes. State the most efficient sorting algorithm and the runtime (giving the tightest bound with no unnecessary constants or lower order terms) to sort the droids by int ID if:

$$k = O(n^2)$$

Algorithm: _____

Runtime: _____

$$k \leq 5$$

Algorithm: _____


Runtime: _____

¹ In case you've forgotten the term "inversion", here's an example that may help: Suppose we have the array [0 1 1 2 3 4 8 6 9 5 7]. This array has 6 inversions: 8-6, 8-5, 8-7, 6-5, 9-5, 9-7.

Login: _____

9. Shortest Paths Algorithms (4.5 Points).

- a. For the graph below, give the order in which Dijkstra's Algorithm would visit each vertex, starting from vertex A.



__A__ ____ ____ ____ ____ ____ ____

- b. Change one of the weights in the graph so that the shortest paths tree returned by Dijkstra's is not correct. Hint: we showed in class that Dijkstra's returns the correct SPT as long as all edges are non-negative.

Set the weight of the edge connecting vertex ____ and vertex ____ to ____.

- c. Suppose we use the following heuristic:

$h(A) = 2$
 $h(B) = 2$
 $h(C) = 20$
 $h(D) = 2$
 $h(E) = 6$
 $h(F) = 2$
 $h(G) = 0$
 $h(H) = 2$

Recall that A* is just Dijkstra's algorithm, except that vertices are given a priority equal to the sum of their Dijkstra priority (distance from the source) plus the heuristic distance, and also that we quit when the target is visited.

Give the path (not order visited) that A* returns from A to G, you may not need all of the blanks:

__A__ ____ ____ ____

10. Hashing and Doubly Linked Lists (7 points). Mu Gulshan, a 61B student, was working on a DLList class. Unfortunately, Mu was transformed into an owl by a jealous Zeus before completing the implementation, so you'll need to pick up where Mu left off. For space reasons, the code for `insertFront` is hidden, but you can assume it works properly. **You may not need all lines.**

```
public class DLList {
    public int size = 0;
    public DNode sentinel = new DNode(null, null, null, null);

    /** Creates a new node and inserts it at the front of the DLL. */
    public void addToFront(String k, String v) {
        insertFront(new DNode(k, v, null, null));
    }

    /** Inserts the given node into the front of the DLL. */
    public void insertFront(DNode toInsert) { ... }

    /** Returns node containing k, or null if node doesn't exist. */
    public DNode getNode(String k) {
        _____
        _____
        _____
        _____
    }

    /** Unlinks the given node from its position in the list. */
    public void detachNode(DNode toDetach) {
        _____
        _____
        _____
        _____
    }

    public static class DNode {
        public String key, value;
        public DNode prev, next;
        public DNode(String k, String v, DNode p, DNode n) {
            key = k;
            value = v;
            prev = p;
            next = n;
        }
    }
}
```

Login: _____

Mu built the strange DLL class listed on the previous page as a support class for an optimized HashMap. The main idea is that recently used items are kept at the front of each list. In Mu's scheme, the HashMap (which maps Strings to Strings), behaves as follows:

1. Key/value pairs are stored in nodes in an array of buckets, where each bucket is a doubly linked list.
2. When a key/value pair is inserted into a bucket, we check to see if the key is already in the bucket.
 - If the key/value pair does not exist in the bucket, we create a new node (with that pair) at the front of the bucket.
 - If the key/value pair exists in a node in the bucket, we update that node's value, and move that node to the front.

Mu's last commit to the DLLList class was the following method:

```
/** Returns node associated with string k or null if doesn't exist.
 * Moves item to the front if it does exist. */
public DNode getNodeAndPromote(String k) {
    DNode p = getNode(k);
    if (p == null) { return null; }
    detachNode(p);
    insertFront(p);
    return p;
}
```

Using this method (and others), fill in the put method below. **You may not need all lines.**

```
public class DLLHashMap {
    public DLLList[] data;
    public DLLHashMap(int numBuckets) {
        data = new DLLList[numBuckets];
        for (int i = 0; i < numBuckets; i += 1)
            data[i] = new DLLList();
    }

    /** If k exists, update value and move to front. Otherwise add new
     * node (containing key and value) to front. */
    public void put(String k, String v) {
        _____
        _____
        _____
        _____
        _____
        _____
        _____
        _____
    }
}
```

11. PNH (0 points). This remote surfing spot was “discovered” via Google Earth in the 2000s, and is known as one of the most magnificently long barreling waves in the world. Either of its common names (as given by its discoverers, or by the people already surfing it before it was discovered) is OK.

12. Radix Sorting and Bits (8 Points).

- a. In LSD radix sorting, we sort our inputs digit by digit, starting from the least significant digit and working our way leftwards. For each digit, we used counting sort. Conceptually, we can imagine that we used a subroutine `countSort(Something[] a, int k)`, which sorts the array `a` of something based on the `k`th digit of each object in `a`.

Suppose that we replaced counting sort with a special version of insertion sort `insertionSort(Something[] a, int k)` that insertion sorts based on only the `k`th digit of each `Something`. Would LSD radix sort still work? What would be the worst-case runtime in terms of `N` and `W`, where `N` is the number of keys and `W` is the width of the keys in digits? Justification is optional but may be considered for partial credit.

Correct (yes/no): _____

Runtime, big O: _____

- b. Same question as a, but using `heapSort(Something[] a, int k)` to sort on each digit instead of counting sort: Would this version of LSD radix sort yield the correct result? What would be the worst case runtime in terms of `N` and `W`? Justification is optional but may be considered for partial credit.

Correct (yes/no): _____

Runtime, big O: _____

- c. The absolute value function in Java can be implemented as:

```
public static int abs(int x) {  
    if (x < 0) { return (~x + 1); }  
    return x;  
}
```

This function works correctly for every integer value in Java except one. Which value is this, and why? Recall that the `~x` operation flips all the bits of `x`.

Login: _____

d. Suppose we are given an LSD radix sorting function described below.

```
public class LSDRadixSorter {
    /** Sorts input digit-by-digit. */
    public static void LSDRadixSort(Digitable[] a) { ... }
}
```

LSDSorter requires that the inputs implement the Digitable interface, described below:

```
public interface Digitable {
    /** Returns the number of digits. */
    public int numDigits();
    /** Returns the kth digit. */
    public int kthDigit(int k);
}
```

Fill in `kthDigit` so that the result of our LSD radix sort is the same as if we used a regular comparison based sort. You may not need all lines given.

```
public class KetchupFriend implements Comparable<KetchupFriend>, Digitable {
    public int ketchupLevel;    // assume non-negative
    public int rednessQuotient; // assume non-negative
    public int compareTo(KetchupFriend other) {
        if (ketchupLevel > other.ketchupLevel) { return 1; }
        else if (ketchupLevel < other.ketchupLevel) { return -1; }
        if (rednessQuotient > other.rednessQuotient) { return 1; }
        else if (rednessQuotient < other.rednessQuotient) { return -1; }
        return 0;
    }

    public int numDigits() {
        return 64;
    }

    /** Returns kth character where k = 0 is the least
     * significant digit. */
    public int kthDigit(int k) {
        _____
        _____
        _____
        _____
        _____
    }
}
```

13. Shortest Paths Algorithm Design (5 points).

Warning: This problem is particularly challenging. **Do not start until you feel like you've done everything else you can.** We will be award very little partial credit for this problem. Solutions which are correct but do not meet our time and space requirements (below) will be not be awarded credit.

- a. Design an efficient algorithm for the following problem: Given a weighted, directed graph G where the weights of every edge in G are all integers between 1 and 10, and a starting vertex s in G , find the distance from s to every other vertex in the graph (where the distance between two vertices is defined as the weight of the shortest path connecting them, or infinity if no such path exists).

Your algorithm must run faster than Dijkstra's to receive credit.

- b. Give the running time of your algorithm in terms of V and E (the number of vertices in the graph and the number of edges in the graph,