

# University of California, Berkeley – College of Engineering

Department of Electrical Engineering and Computer Sciences

Spring 2014

Instructor: Dr. Dan Garcia

2014-03-12



## CS61C Midterm



After the exam, indicate on the line above where you fall in the emotion spectrum between “sad” & “smiley”...

Last Name	<b>Answer Key</b>
First Name	
Student ID Number	
Login	cs61c-
Login First Letter (please circle)	a b c d e f g h i j k l m n o p q r s t u v w x y z
Login Second Letter (please circle)	a b c d e f g h i j k l m n o p q r s t u v w x y z
The name of your <b>LAB</b> TA (please circle)	Alan   Jeffrey   Kevin   Roger   Sagar   Shreyas   Sung Roa   William
Name of the person to your Left	
Name of the person to your Right	
All the work is my own. I had no prior knowledge of the exam contents nor will I share the contents with others in CS61C who have not taken it yet. (please sign)	

### Instructions (Read Me!)

- Don't Panic!
- This booklet contains 6 numbered pages including the cover page. Put all answers on these pages; don't hand in any stray pieces of paper.
- Please turn off all pagers, cell phones & beepers. Remove all hats & headphones. Place your backpacks, laptops and jackets at the front. Sit in every other seat. Nothing may be placed in the “no fly zone” spare seat/desk between students. No computing devices allowed.

Question	1	2	3	4	Total
Minutes	48	24	24	24	120
Points	30	15	15	15	75
Score	30	15	15	15	75

**Question 1: What's that funky smell?! Oh yeah, it's potpourri...** (48 min, 30 pts)

- a) The *unsigned distance* between two bit patterns is the absolute value of the difference of their values, interpreted as unsigned numbers. Rank the following according to the unsigned distance between -1 and 0 (+0 if a representation has multiple zeros) in that representation. You should assign a rank of 1 to the representation with the ***smallest*** unsigned distance between -1 and 0.

<code>int64_t</code>	__ <b>5</b> __
64-bit One's Complement	__ <b>4</b> __
64-bit Sign and Magnitude	__ <b>2</b> __
64-bit Bias notation	__ <b>1</b> __
<code>double</code>	__ <b>3</b> __

Show all your work here!

- b) As defined in IEEE 754-2008 standard, *half-precision floating point* (FP) is a 16-bit FP representation: 1 sign bit, 5 exponent bits, and 10 significand bits. The exponent bias of 15. What is the binary representation of the ***smallest half-precision float which is strictly larger than 1***? What is its value? Leave your answer in terms of powers of two.

0b **0 01111 0000000001** =  **$1 + 2^{-10}$**

Show all your work here!

- c) How would J-type instructions be affected (in terms of their “reach”) if we relaxed the requirement that instructions be placed on word boundaries, and instead required them to be placed on *half-word* boundaries.

**The range over which we can jump would be cut in half.**

- d) Building on the idea from the previous question, give a minor tweak to the MIPS ISA to allow us to use *true absolute addressing* (i.e., maximal “reach”) for all J-type instructions.

**Only allow jumps to addresses which are multiples of  $2^6$ .**

- e) Assume a request for 100 Bytes on either the **Stack** and **Heap** would succeed. In the worst case, *how many clock cycles* would it take to **allocate** 100 Bytes in each area & *why*? (“How many” should be answered with one of: {1, tens, thousands+})

**Stack:** **1** because **it's just `addiu $sp $sp -100`**  
**Thousands+**  
**Heap:** **the freelist might contain lots of slivers to check**

**Question 1: *What's that funky smell?! Oh yeah, it's potpourri... (continued)*** (48 min, 30 pts)

- f) You have a program that can achieve almost a 20x speedup with millions of processors, so what is the percent of the parallel portion of its code?

\_\_\_ **95** \_\_\_

*Show all your work here!*

- g) Suppose the assembler knew the file line numbers of all labels before it began its first pass over a file, and that every line in the file contains an instruction. Then the assembler would need \_\_\_ **2** \_\_\_ pass(es) to translate a MAL file, and \_\_\_ **1** \_\_\_ pass(es) to translate a TAL file. These numbers differ because of \_\_\_ **pseudoinstructions** \_\_\_ (write n/a if they don't differ).

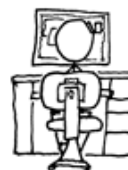
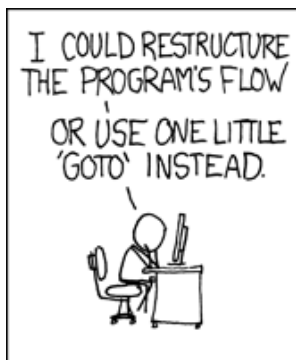
- h) Complete the code below, using *only one TAL instruction*, so that it returns true iff \$a0 is an I-type instruction or a J-type instruction, and then translate the instruction into binary then hexadecimal.

**srli \$v0, \$a0, 26      0000|00 00|000 0|0100 0001|0 110|10 00|0010    0004 1682**  
IJ-instr: \_\_\_\_\_ → 0b \_\_\_\_\_ → 0x \_\_\_\_\_  
          jr \$ra

*Show all your work here!*

- i) What is one thing Google did to increase their Power Usage Efficiency (PUE)?

**Careful airflow handling, elevated cold aisle temp, use free cooling, per-server ups**



## Question 2: *Running in circles* (24 min, 15 pts)

- a) Recall the exercise `ll_cycle` from lab 2, in which we checked if a linked list contained a cycle using the *tortoise and hare* algorithm. We've provided you with most of a simple *recursive* implementation in the space below. Fill in the base cases.

```
int ll_has_cycle(node *ptr) {
    if (!ptr)
        return 0;
    return has_cycle(ptr, ptr->next);
}
```

```
typedef struct node {
    int value;
    struct node *next;
} node;
```

```
int has_cycle(node *tortoise, node *hare) {
    hare == tortoise
    if (_____)
        return 1;
    !hare || !hare->next
    if (_____)
        return 0;

    return has_cycle(tortoise->next, hare->next->next)
}
```

- b) Now that you've warmed up on the C version of this code, let's convert `has_cycle` into **recursive** MAL MIPS. Assume that the fields of the `structs` are not permuted from the `struct` definition. You may use fewer lines than we provide you, but do not add any more than the space provided.

# \$a0 contains the pointer to the tortoise, \$a1 contains the pointer to the hare.

has\_cycle: li \$v0 1

beq \$a0 \$a1 done

li \$v0 0

\$a1 \$0

beq \_\_\_\_\_ done

lw \$a1 4(\$a1)

\$a1 \$0

beq \_\_\_\_\_ done

lw \$a0 4(\$a0)

lw \$a1 4(\$a1)

\$sp \$sp -4

addiu \_\_\_\_\_

sw \$ra 0(\$sp)

jal has\_cycle

lw \$ra 0(\$sp)

\$sp \$sp 4

addiu \_\_\_\_\_

done: jr \$ra

← circled, changed to nop

← circled, changed to nop

← change to j has\_cycle

- c) You want to change this to be an *iterative* MIPS solution, but you want to change the fewest lines you can. Circle those lines you would change to make the program work *iteratively*.

**Question 3: *Our band is called 1023MiB... We haven't had any gigs yet.*** (24 min, 15 pts)

We have a standard 32-bit byte-addressed MIPS machine with a 1KiB direct-mapped write-back cache and 16B block size.

a) How many bits are used for the Tag? 22 ...Index? 6 ...Offset? 4

Consider the following C code, and answer the questions below. `a` and `s` are pointers to 8-bit unsigned integer arrays, of the same size (a multiple of the cache size) that are aligned on 16-byte boundaries. The arrays contain only one `0x00`, in the last byte. `a` and `s` are not necessarily distinct.

```
void our_strcpy(uint8_t *d, uint8_t *s) {
    char c;
    do {
        c = *s;
        *d = c;
        s++; d++;
    } while (c);
}
```

b) What is the *lowest* possible cache hit rate for `our_strcpy`? 0

**Compulsory and conflict**

c) What *types* of misses are there? \_\_\_\_\_

d) What is the *smallest possible value* of  $(a - s)$  that would get this hit rate? 1 KiB

e) What is the *highest* possible cache hit rate for `our_strcpy`? 31/32

f) What is one possible value of  $(a - s)$  where we would get this hit rate? 0

**2 misses per block \* 2^6 blocks/cache \* 2^13 caches = 2^20 misses**

g) If we ran `our_strcpy` with a 4-way set-associative LRU cache, and the size of both `a` and `s` is 8MiB, what is the *most # of misses* possible? 1 Mebi

Show all your work here!

#### Question 4: A bad case of Not Invented Here Syndrome... (24 min, 15 pts)

- a) A colleague of yours has implemented some homebrew C99 string manipulation functions, while steadfastly refusing to use any standard libraries, but they're buggy! We've marked each potentially problematic line with `// <number>`. Your job is to fill in a correct replacement line in the corresponding row of the following table, or write 'OK' if there is nothing wrong. DO NOT LEAVE ANY FIELDS BLANK, or we will assume you just didn't get to this part of the exam.

Line number	Replacement Code
1	<code>for(char c = *s; (c = *s) != '\0'; s++) {</code>
2	<code>*s += 'a' - 'A';</code>
3	<code>OK</code>
4	<code>OK</code>
5	<code>odds += numbers[i] &amp; 1    ...    odds += *(numbers+i) &amp; 1</code>

*/\*\* Converts the string S to lowercase \*/*

```
void string_to_lowercase(char *s) {  
    for(char c = *s; c != '\0'; s++) {           // 1  
        if(c >= 'A' && c <= 'Z') {               // 2  
            s += 'a' - 'A';  
        }  
    }  
}
```

*/\*\* Returns the number of bytes in S before, but not counting, the null terminator. \*/*

```
size_t string_length(char *s) {  
    char *s2 = s;  
    while(*s2++);                               // 3  
    return s2 - s - 1;                           // 4  
}
```

*/\*\* Return the number of odd numbers in a number array \*/*

```
uint32_t number_odds(uint32_t *numbers, uint32_t size) {  
    uint32_t odds = 0;  
    for (uint32_t i = 0; i < size; i++)  
        odds += *numbers+i && 1;                // 5  
    return odds;  
}
```

