Computer Science 162, Fall 2014
David Culler
University of California, Berkeley
Midterm 1
September 29, 2014

Name	
SID	
Login	
TA Name	
Section Time	

This is a closed book exam with one 2-sided page of notes permitted. It is intended to be a 50 minute exam. You have 80 minutes to complete it. The number at the beginning of each question indicates the points for that question. Write all of your answers directly on this paper. Make your answers as concise as possible. If there is something in a question that you believe is open to interpretation, please raise your hand to request clarification. When told to open the exam, put your login on every page and check that you have them all. (Final page is for reference.)

By my signature below, I swear that this exam is my own work. I have not obtained answers or partial answers from anyone. Furthermore, if I am taking the exam early, I promise not to discuss it with anyone prior to completion of the regular exam, and otherwise I have not discussed it with anyone who took the early alternate exam.

X _____

Grade Table (for instructor use only)

Question	Points	Score
1	25	
2	25	
3	25	
4	25	
Total:	100	

1. (25 points) **Operating Systems Concepts**

   (a) (20 points) Choose **either** true or false for the below questions. You do not need to provide justifications.

       i. (2 points) A user program can request an operating system service by issuing a system call.

       ii. (2 points) Threads in a single process are prevented by the operating system from accessing the stacks of other threads.

       iii. (2 points) A pintos kernel thread enters `thread_switch` with interrupts disabled, so `thread_switch` enables them before switching to the next thread.

       iv. (2 points) The scheduler places RUNNING threads on a queue according to a scheduling algorithm.

       v. (2 points) In operating systems like Unix, different operations are used for accessing files, transferring information through sockets, and communication between processes.

       vi. (2 points) In operating systems like Unix, different operations are used for setting file transfers and network communication.

       vii. (2 points) Interrupt handlers can call `pthread_mutex_lock` because it is atomic.

       viii. (2 points) Disabling interrupts always protects a critical section of code.

       ix. (2 points) A successful call to fork() always returns twice - once in the child process and once in the parent process.

       x. (2 points) Path resolution is implemented by the kernel.

(b) (5 points) Which of the following is true about Round Robin Scheduling? Select **all** the choices that apply.

Cache performance is likely to improve relative to FCFS.

If no new threads are entering the system all threads will get a chance to run in the cpu every `QUANTA*SECONDS_PER_TICK*NUMTHREADS` seconds. (Assuming `QUANTA` is in ticks).

This is the default scheduler in Pintos

If quanta is constantly updated to become the # of cpu ticks since boot, Round Robin becomes FIFO.

If all threads in the system have the same priority, Priority Schedulers **must** behave like round robin.

2. (25 points) **Processes and Threads**
For the following consider the following nearly identical process based and thread based programs. For each question, answer YES or NO and provide a short, crisp explanation. **Assume that functions will be well-behaved (e.g., malloc will succeed) and that printf() is atomic.** (We will answer half the first part to provide an example.)

```
1   int i = 100;
2   char *buf;
3
4   void process() {
5     pid_t pid;
6     int status, rd;
7     int j = 1;
8     buf = strcpy(malloc(100), "boring");
9     pid = fork();
10    if (pid != 0) {                    /* parent */
11      printf("P␣Parent:␣j=%d\n", j);
12      i = 162;
13      j = 2;
14      printf("P␣Parent:␣j=%d\n", j);
15      printf("P␣Parent:␣%s\n", buf);
16      j = 4;
17      wait(&status);
18    } else {                           /* child */
19      printf("P␣Child:␣i=%d,␣j=%d\n", i, j);
20      printf("P␣Child:␣i=%d,␣j=%d\n", i, j);
21      j = 3;
22      strcpy(buf, "cool");
23      exit(0);
24    }
25  }
```

```
1   int i = 100;
2   char *buf;
3
4   void *tfun(void *noarg) {
5     int j = 0;
6     printf("T␣Child:␣i=%d,␣j=%d\n", i, j);
7     printf("T␣Child:␣i=%d,␣j=%d\n", i, j);
8     j = 3;
9     strcpy(buf, "cool");
10    pthread_exit(NULL);
11  }
12
13  void thread() {
14    pthread_t tid;
15    int j = 1;
16    buf = strcpy(malloc(100), "boring");
17    pthread_create(&tid, NULL, tfun, NULL);
18    printf("T␣Parent:␣j=%d\n", j);
19    i = 162;
20    j = 2;
21    printf("T␣Parent:␣j=%d\n", j);
22    printf("T␣Parent:␣%s\n", buf);
23    j = 4;
24    pthread_join(tid, NULL);
25  }
```

(a) (5 points) Can the output of the process-based program ever include (meaning they appear in this order, possibly with other output interleaved) the following? Why? (We are going to answer this one for you!)

```
P Parent: j=1
P Child: i=100, j=1
P Parent: j=2
```

Can the output of the thread-based program ever include the following? Why?

```
T Parent: j=1
T Child: i=100, j=1
T Parent: j=2
```

(b) (10 points) (be careful) Can the output of the process-based program ever include the following? Why?

```
P Parent: j=1
P Parent: j=2
P Child: i=100, j=1
```

Can the output of the thread-based program ever include the following? Why?

```
T Parent: j=1
T Parent: j=2
T Child: i=100, j=1
```

(c) (10 points) Can the output of the process-based program ever include the following? Why?

```
P Parent: cool
```

Can the output of the thread-based program ever include the following? Why?

`T Parent: cool`

3. (25 points) **Synchronized Objects**

(a) (15 points) Assuming malloc is threadsafe, modify (by showing where code needs to be inserted) the following unbounded stack abstraction to make **ppush** and **ppop** threadsafe using pthreads operations. We have started it for you by adding the lock to the pdl struct. (You do not need to eliminate the busy-wait. It only needs to be threadsafe, not highly efficient. You may not need to fill every line indicated.)

(b) (10 points) Fill in **dequeue** to create a thread-safe function that waits until the stack is non-empty, removes and item and returns it.

```
typedef struct item {
  struct item *next;
  void *val;
} item_t;

typedef struct pdl {
  item_t *head;

  --------------------------
} pdl_t;

pdl_t *new_pdl() {
  pdl_t *s = malloc(sizeof(pdl_t));
  s->head = NULL;
  pthread_mutex_init(&s->pdl_lock, NULL);
  return s;
}

void ppush (pdl_t *s, void *val) {

  -------------------------------
  item_t *new_item;

  ----------------------------
  new_item = malloc(sizeof(item_t));
  new_item->val = val;
  new_item->next = s->head;
  s->head = new_item;

  -----------------------------
}
```

```
void *ppop(pdl_t *s) {



  ----------------------------
  item_t *pop;
  pop = s->head;
  void *res = NULL;
  if (pop) {

    ------------------------------

    res = pop->val;

    s->head = pop->next;
    free(pop);
  }

  ----------------------------
  return res;
}

bool empty_pdl(pdl_t *s) {
  return (s->head == NULL);
}

void *dequeue(pdl_t *s) {
  void * res;

  ----------------------------

  ----------------------------

    pthread_yield();
  ----------------------------

  ----------------------------

  return res;
}
```

4. (25 points) **Operating System Implementation: Pintos**
   You design a new scheduler, you call it TFS. The idea is relatively simple, if a thread makes it
   to the end of a round robin scheduling quanta it sets its own priority to -1 times the number of
   ticks it has **ever** spent in the cpu.

   You may make the following assumptions in this problem.

   - Priority scheduling in Pintos is functioning properly,

   - Priority donation is not implemented.

   - Alarm is not implemented.

   - `thread_set_priority` is never called by the thread

   - You may ignore the limited set of priorities enforced by pintos (priority values may span
     from `INT_MIN` to `INT_MAX`).

   i. (4 points) Below is the declaration of `struct thread`, what field(s) would we need to add
      to make TFS possible. You may not need all the blanks.

```
struct thread
  {
    /* Owned by thread.c. */
    tid_t tid;                          /* Thread identifier. */
    enum thread_status status;          /* Thread state. */
    char name[16];                      /* Name (for debugging purposes). */
    uint8_t *stack;                     /* Saved stack pointer. */
    int priority;                       /* Priority. */
    struct list_elem allelem;           /* List element for all threads list. */


    /* Shared between thread.c and synch.c. */
    struct list_elem elem;              /* List element. */

#ifdef USERPROG
    /* Owned by userprog/process.c. */
    uint32_t *pagedir;                  /* Page directory. */
#endif

    _____          /* What goes here? */

    _____          /* What goes here? */

    _____          /* What goes here? */

    /* Owned by thread.c. */
    unsigned magic;                     /* Detects stack overflow. */
  };
```

i

ii. (8 points) A what is needed for `thread_tick()` for TFS to work properly. You may not need all the blanks.

```
void
thread_tick (void)
{
  struct thread *t = thread_current ();

  /* Update statistics. */
  if (t == idle_thread)
    idle_ticks++;
#ifdef USERPROG
  else if (t->pagedir != NULL)
    user_ticks++;
#endif
  else
    kernel_ticks++;


  _____;


  _____;
  /* Enforce preemption. */
  if (++thread_ticks >= TIME_SLICE) {
    intr_yield_on_return ();

    _____;

  }



}
```

iii. (5 points) In a few short words, describe the overall behavior of this scheduler. How is it different from regular round robin or FIFO scheduler.

iv. (4 points) What could you do to "trick" this scheduler to giving a thread more cpu time. (Hint: `thread_ticks` is set to 0 every time the scheduler picks a new thread)

v. (4 points) What is a flaw in this scheduler when you have a steady state of threads being created, and very little threads destroyed. (Hint: Think about what happens when the scheduler has run for a really long time).

```
/****************************** Pthreads ******************************/
int pthread_create(pthread_t *thread, const pthread_attr_t *attr,
                          void *(*start_routine) (void *), void *arg);
int pthread_join(pthread_t thread, void **retval);
int pthread_mutex_init(pthread_mutex_t *restrict mutex,
    const pthread_mutexattr_t *restrict attr);
int pthread_mutex_lock(pthread_mutex_t *mutex);
int pthread_mutex_trylock(pthread_mutex_t *mutex);
int pthread_mutex_unlock(pthread_mutex_t *mutex);
int sem_init(sem_t *sem, int pshared, unsigned int value);
int sem_post(sem_t *sem);
int sem_wait(sem_t *sem);
int sem_trywait(sem_t *sem);

/**************************** Strings & Processes ****************************/
char *strcpy(char *dest, const char *src);
char *strncpy(char *dest, const char *src, size_t n);

pid_t wait(int *status);
pid_t waitpid(pid_t pid, int *status, int options);


/**************************** Pintos   ****************************/
```